

SuccessApp Memory Leak Incident

SuccessApp is one of many of our API's currently that is feeding data for reporting into our Heroku Environment. Per the API, there are new fields that are being added on a time basis, and some fields contain nested JSON which take a lot of resources when flattening in order to fit our model.

Example of a Nested JSON Structure from SuccessApp:

```
"territories": [  
  {  
    "id": "04b191f8-c86b-46fd-9932-23a592562349",  
    "name": "LATAM",  
    "color": "indigo-dark",  
    "created_at": "2023-04-23 21:00:18",  
    "updated_at": "2023-04-23 21:00:18",  
    "association": {  
      "source_id": "ea6338ef-fe1c-4aa8-9e56-99d24e72bb7a",  
      "target_id": "04b191f8-c86b-46fd-9932-23a592562349",  
      "id": "5e1e6be8-b9c5-4665-a25d-31baccb2f31d"  
    }  
  },  
  {  
    "id": "f79e74c6-1253-473a-90dc-cb3ff5dddc2f7",  
    "name": "LATAM-SOLA-BRAZIL SOUTH CROSS",  
    "color": "indigo-light",  
    "created_at": "2023-04-23 21:00:20",  
    "updated_at": "2023-04-23 21:00:20",  
    "association": {  
      "source_id": "ea6338ef-fe1c-4aa8-9e56-99d24e72bb7a",  
      "target_id": "f79e74c6-1253-473a-90dc-cb3ff5dddc2f7",  
      "id": "590b85a8-a136-4ae9-ac98-4d406d6614ab"  
    }  
  }  
]
```

Notice how there are multiple fields under the parent Territories field. On the reporting side, we can't use a nested JSON so flattening it out into it's separate column is the requirement.

This is just one of many parent fields that SuccessApp contained, but below is an example of how territories is flattened inside Heroku.

Heroku Flattened out Column Structure:

ABC territories_id	▼	ABC territories_color	▼	ABC territories_created_at	▼	ABC territories_updated_at	▼
ABC territories_association.target_id				▼	ABC territories_association.id		
					▼		

These nested fields are very expensive in memory especially when merging everything into a single data frame so optimization is essential when sending into Heroku. Currently our Heroku environment is on a private-L dyno which supports up to **14GB** of ram.

Optimization:

SuccessApp also contains many duplicate records and after joining each parent field dataset together and using the final one to write to Postgres, it totaled up to **135GB** of ram. The solution to optimizing such a huge dataset is to drop duplicate records **BEFORE** each merge.

```

1 df_final = ft.reduce(lambda left, right: pd.merge(left, right, on=['id'], how='outer'), dfs)
2 df_final = df_final.drop('bridges', axis=1)
3 df_final = df_final.drop('territories', axis=1)
4 df_final = df_final.drop('tags', axis=1)
5 df_final = df_final.drop('stage_history', axis=1)
6 df_result = df_result.drop_duplicates()
7 df_final_r = pd.merge(df_final, df_result, on='result_id', how='left')

8 df_final_r = df_final_r.drop('result', axis=1)
9 df_final_r['id_p'] = df_final_r['id']
0 df_final_r = df_final_r.drop('id', axis=1)
1 df_result_reasons = pd.concat([pd.DataFrame(x) for x in df_final_r['result_reasons']], keys=df_final_r['id_p']).reset_index()
2 df_result_reasons['result_reasons_id'] = df_result_reasons['id']
3 df_result_reasons = df_result_reasons.drop('id', axis=1)
4 df_result_reasons['id'] = df_result_reasons['id_p']
5 df_result_reasons = df_result_reasons.drop('id_p', axis=1)
6 df_result_reasons['result_reasons_title'] = df_result_reasons['title']
7 df_result_reasons['result_reasons_created_at'] = df_result_reasons['created_at']
8 df_result_reasons['result_reasons_updated_at'] = df_result_reasons['updated_at']
9 df_result_reasons = df_result_reasons.drop('title', axis=1)
0 df_result_reasons = df_result_reasons.drop('created_at', axis=1)
1 df_result_reasons = df_result_reasons.drop('updated_at', axis=1)
2 df_final_r['id'] = df_final_r['id_p']
3 df_final_r = df_final_r.drop('id_p', axis=1)
4 df_final_r = df_final_r.drop('result_reasons', axis=1)
5 df_final_r = df_final_r.drop_duplicates()
6 df_result_reasons = df_result_reasons.drop_duplicates()
7 df_final_update1 = pd.merge(df_final_r, df_result_reasons, on='id', how='left')
8 df_final_update1['date_inserted'] = datetime.now()

```

```
1 df_final_update1.memory_usage(index=True).sum()
```

161600536

The final merged dataset we get totals to just 161600536 bytes which is only just 0.161600536GB of ram. Dropping duplicates before a huge merge can be beneficial and should be considered whenever large datasets are involved!